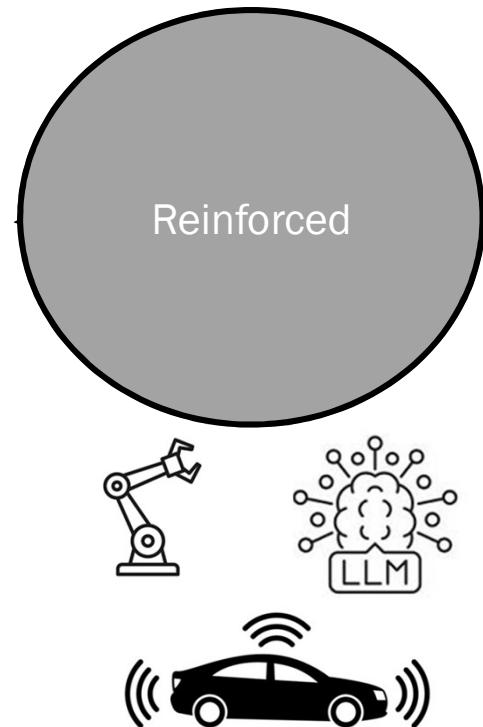
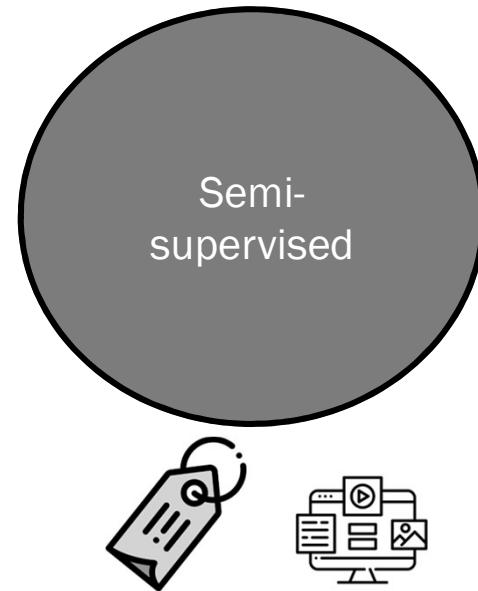
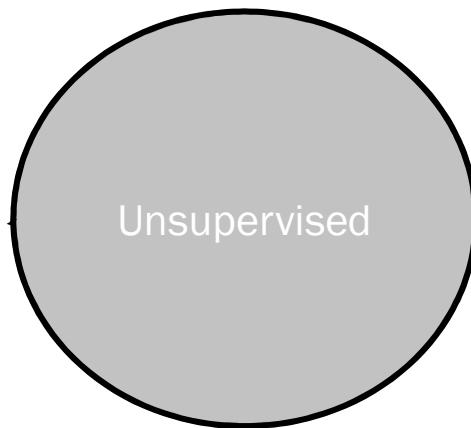
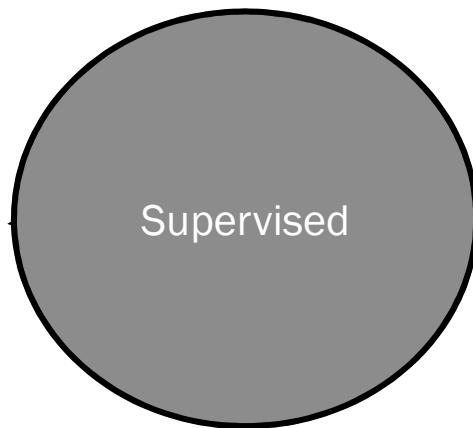


ML-3

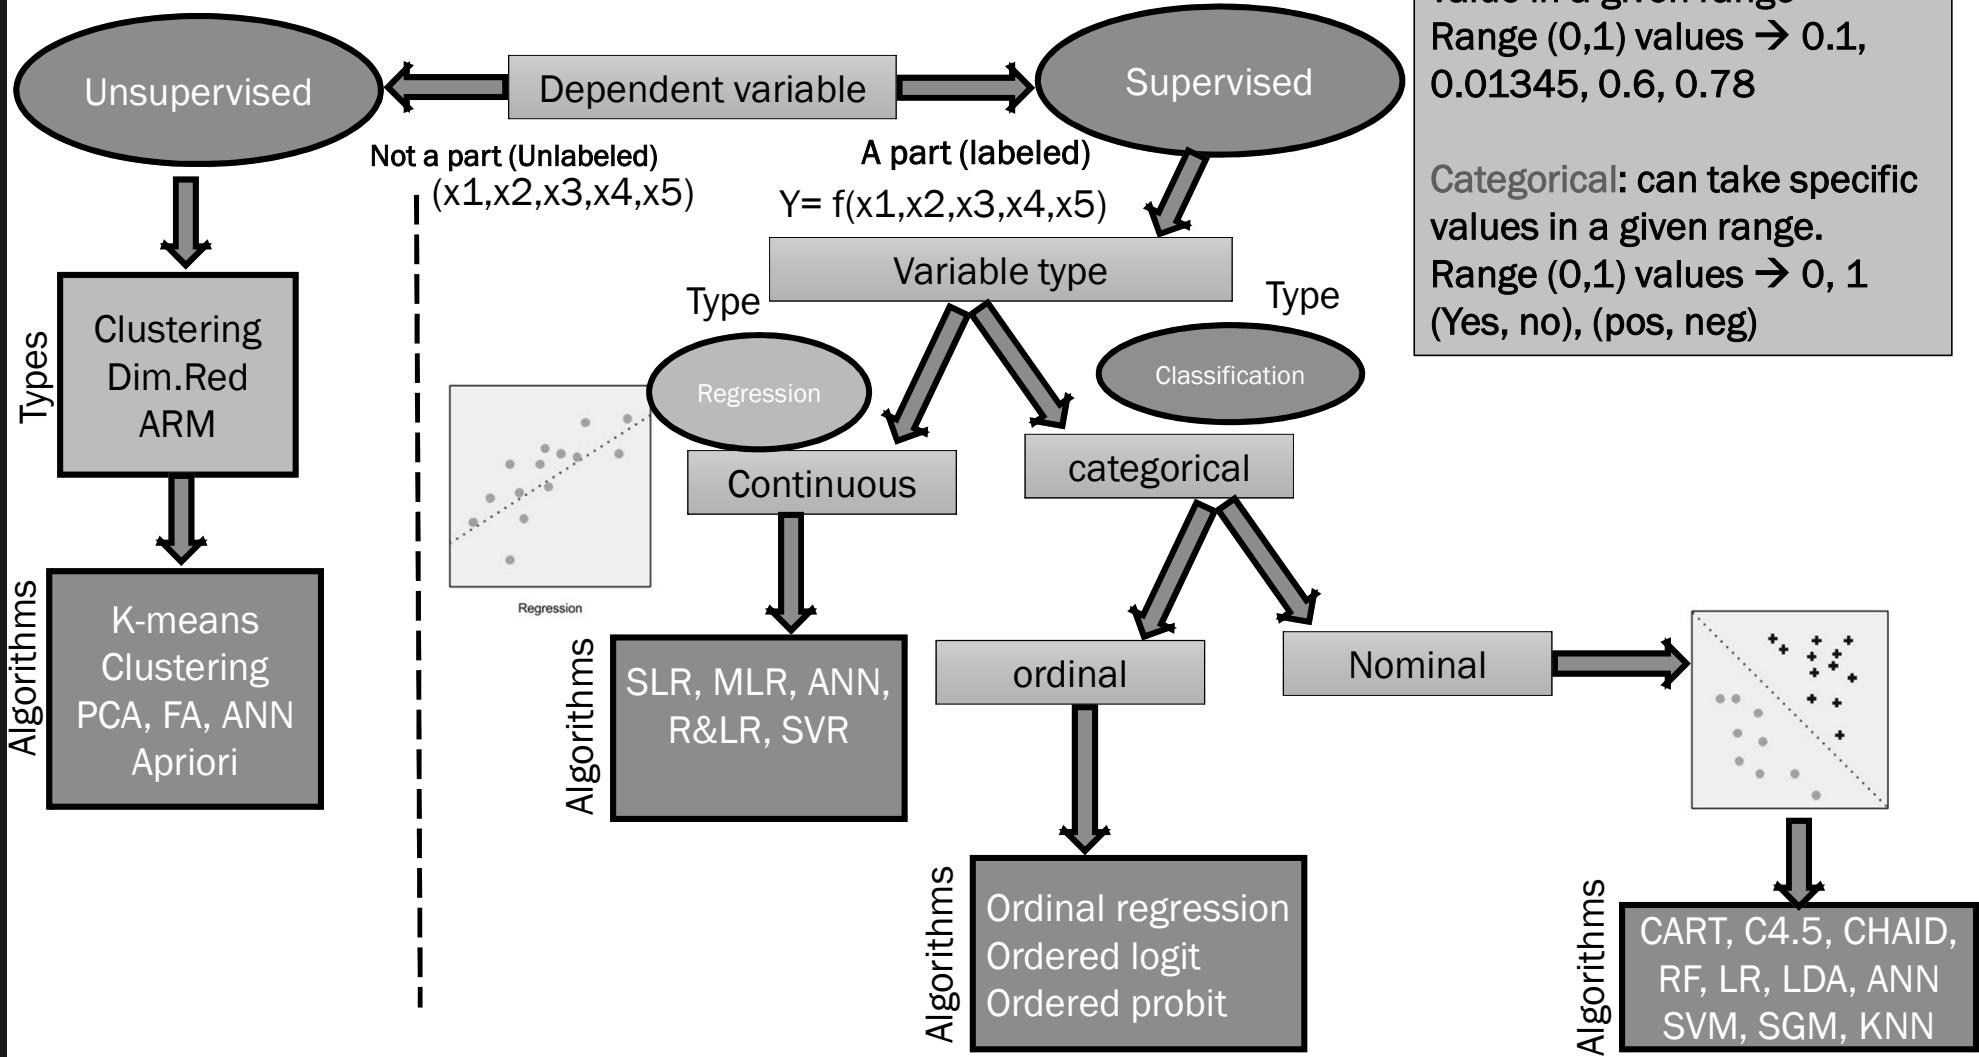
What's covered

- ML, ML types
- examples
- Few jargon terms
- Unsupervised learning

ML Types



Machine learning



Supervised Vs. Unsupervised

Suppose you are given the variables/ features (Cloud density, Relative Humidity, Pressure, Temperature, wind speed and direction, rain fall, rain fall metric) that indicate the weather in different cities (Bangalore, Hyderabad, Chennai, Pune, Mysore.....)

City	D	RH	P	T	speed	direction	Rain fall	RF metric
Bangalore								
Mumbai								
Delhi								
Hyderabad								
Mysore								
Trivandrum								
Pune								
Noida								

Understand the problem
(business/process)
What is the question that
needs to be answered

C	D	RH	P	T	Sp	Direc
Bang						
Mumbai						

$$P^* = a_1 P + a_2 T$$

supervised
Predict a value or class

Unsupervised
Identify pattern

Supervised

Question: will it rain? → classification (given other variables)

Question: how much will it rain? Regression

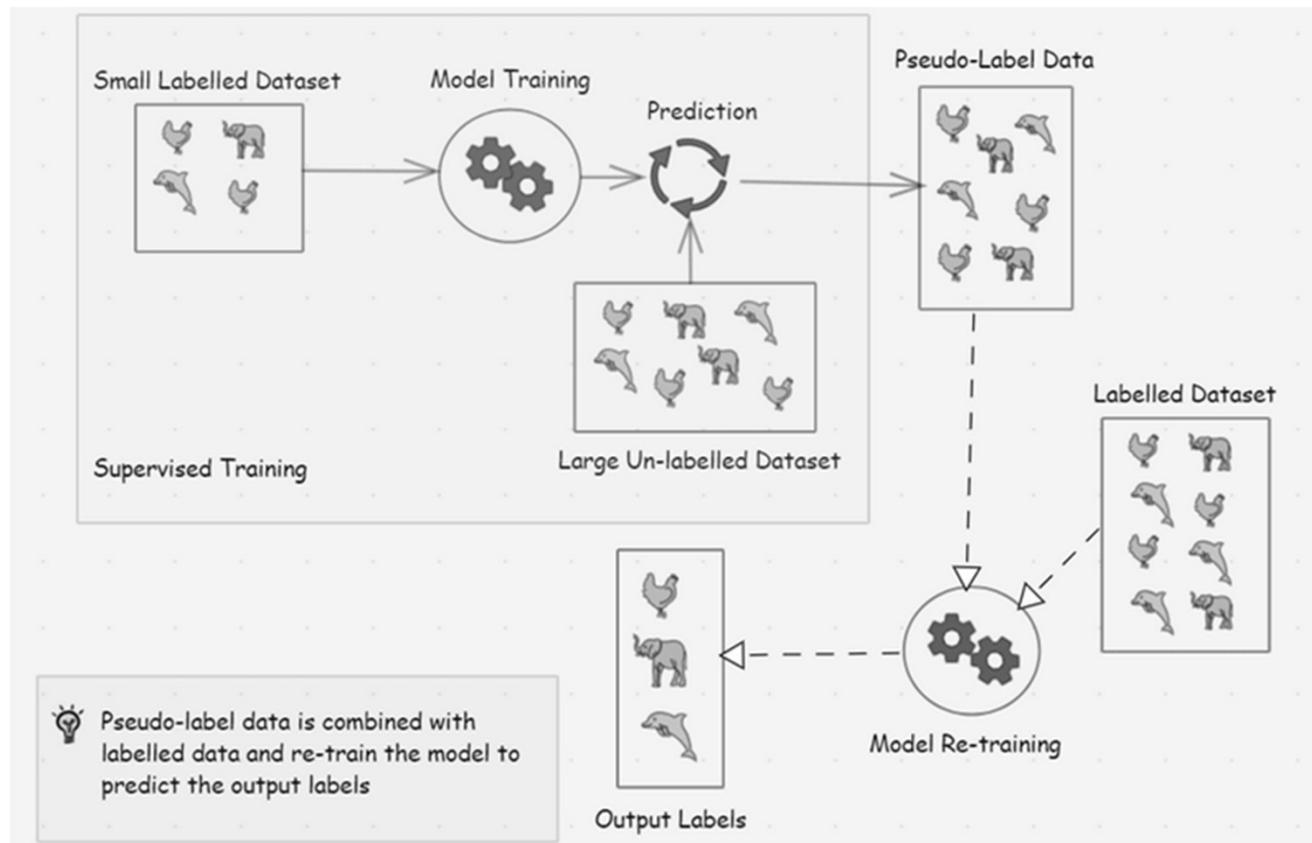
Unsupervised

Question: Which cities are more alike → clustering

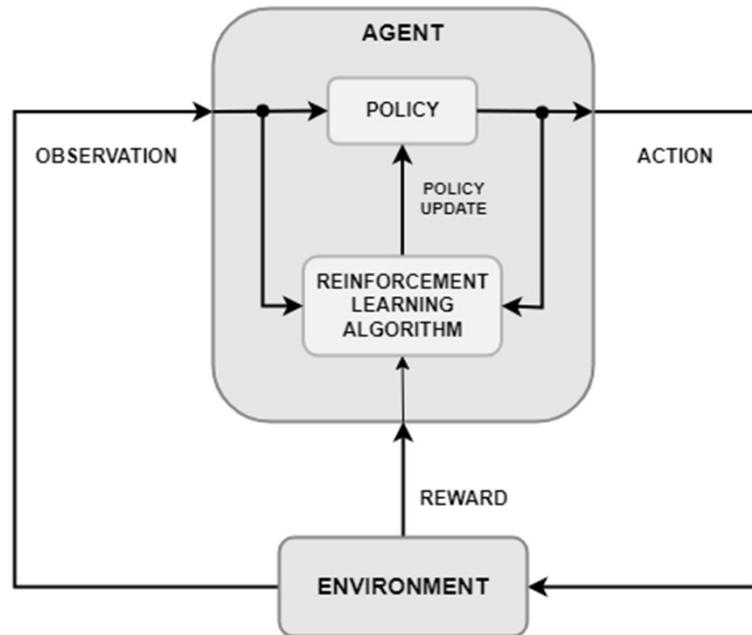
Question: do I need all the six variables to understand the weather → Dimensionality reduction

Question: When the RH is high in Bangalore, is the RH high in Hyderabad too? → Association

Semi-supervised work-flow



Reinforcement learning



Telecommunication company

➤ Scenarios:

- A telecommunication company approaches you and asks you to help them to win their customer loyalty.



Telecommunication company

➤ Scenarios:

- A telecommunication company approaches you and asks you to help them to win their customer loyalty.
- Identify different types of services offered (Data plan, international calls, messaging)
- Identify the customers using similar services to a large extent (cluster)
- Sell custom made services to the target group



Real world applications of clustering

- Document clustering (Ex: Googles news: clustering of news articles based on content)
- location clustering (Ex: Uber: clustering based on popular location)

What's clustering

- Group data such that within group variance < between group variance
- Group such that objects in the same group are more similar to each other in some sense than to objects of different groups.
- These groups are known as clusters and each cluster gets distinct label called cluster ID and the centroid of cluster.
- Clustering types:
 - **Centroid based(ex: K-Means)**
 - **Hierarchical clustering (ex: Agglomerative)**
 - **Density based (ex: DBSCAN)**

Kmeans: Distance formulations

- In KNN method, we calculate the gap/ distance between the new record and the existing class points
- Decide on the number of neighbors to compare with
- Depending on the class of the dominant closest neighbors, decide on the class of the new record
- Distance/ gap calculation

➤ Euclidian($p=2$)

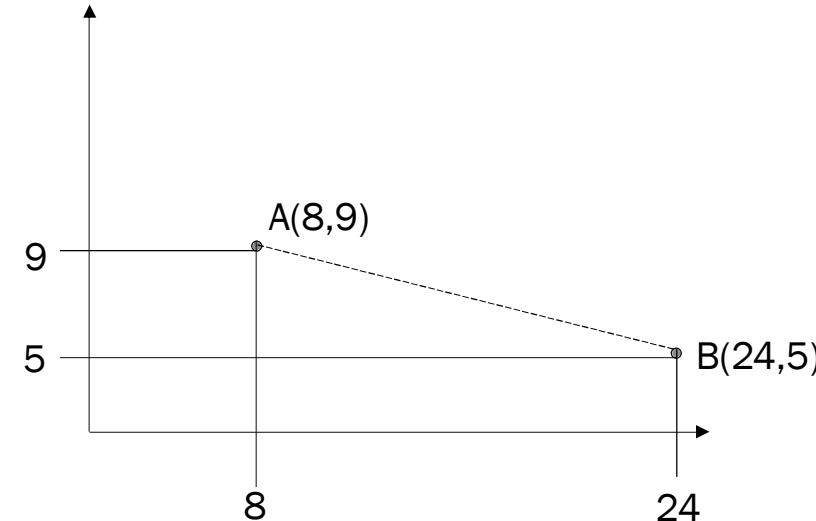
The Minkowski distance of order p between two points

➤ Manhattan($p=1$)

$X = (x_1, x_2, \dots, x_n)$ and $Y = (y_1, y_2, \dots, y_n) \in \mathbb{R}^n$

is defined as:

$$D(X, Y) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$



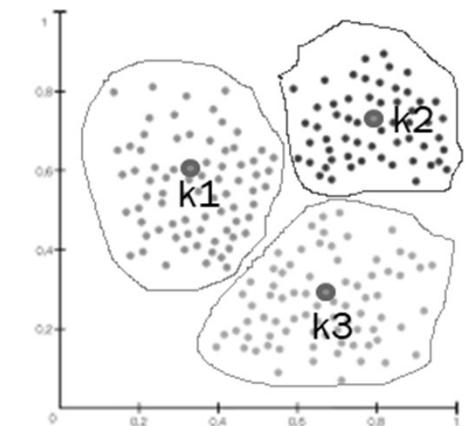
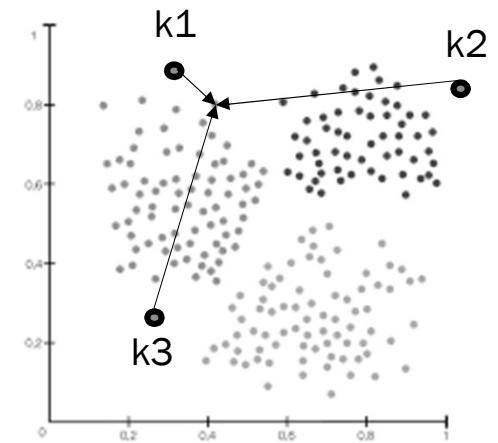
$$E \rightarrow D = \sqrt{(24-8)^2 + (5-9)^2}$$

$$M \rightarrow D = (9-5) + (24-8)$$

$$C \rightarrow D = \max(|9-5|, |24-8|)$$

K-means

- need to specify the number of clusters-K
- Step-1: create arbitrarily ‘k’ number of centroids, assign labels k1,k2,...
- Step-2: calculate the distance of each of the ‘n’ points from each of the ‘k’ centroids
- Step-3: assign each point to nearest centroid (based on distance)
- Step-4: calculate the new centroids based on the records belonging to k_i
- Step-5: check did the cluster index of the points change or iterations less than max: yes → go to step-2. no → stop



Different distance metrics used

Data Type	Metric	Formula	Best Use Case
Numerical	Euclidean	$d(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$	Dense numerical; K-means
Numerical	Manhattan (L1)	$d(p, q) = \sum_{i=1}^n \ p_i - q_i\ $	Sparse/outlier-sensitive
Numerical	Minkowski	$d(p, q) = \left(\sum_{i=1}^n \ p_i - q_i\ ^r \right)^{1/r}$	Tunable ($r=1$ Manhattan, $r=2$ Euclidean)
Categorical	Hamming	$d(p, q) = \sum_{i=1}^n \mathbb{I}(p_i \neq q_i)$	Count mismatches
Mixed	Gower	$d(p, q) = \frac{1}{n} \sum_{i=1}^n \delta_{ij} d_{ij}$	Numerical + categorical

Hamming distance

ID	Color	Size	Shape
1	Red	Small	Circle
2	Red	Small	Square
3	Blue	Small	Circle
4	Blue	Large	Circle
5	Blue	Large	Square

C1 → (Red, Small, Circle)
 C2 → (Blue, Large, Square)

We choose:

- k = 2 clusters
- Distance = Hamming distance

Distance Table

ID	Point	D(C ₁)	D(C ₂)	Assigned
1	(Red, Small, Circle)	0	3	C ₁
2	(Red, Small, Square)	1	2	C ₁
3	(Blue, Small, Circle)	1	2	C ₁
4	(Blue, Large, Circle)	2	1	C ₂
5	(Blue, Large, Square)	3	0	C ₂

Gower distance

Age	Color	HasLoan	
25	Red	Yes	← reference
40	Blue	No	

Feature ranges:

- Age: min=20, max=60

$$|25 - 40| / (60 - 20) = 15 / 40 = 0.375$$

Color

Red ≠ Blue → 1

HasLoan

Yes ≠ No → 1

Gower Distance

$$d = (0.375 + 1 + 1) / 3 = 0.792$$

Different centroid based algorithms

Algorithm	Best For	Centroid Type	Distance Metric	Outlier Robustness	Computational Speed
K-means	Pure numerical data	Arithmetic mean	Euclidean (squared)	Low	Fastest
K-medoids	Numerical, noisy data, custom metrics	Actual data point (medoid)	Any (e.g., Manhattan, Gower)	High	Slow ($O(Kn^2)$)
K-modes	Pure categorical data	Mode (most frequent value)	Simple/Hamming matching	Medium	Fast
K-prototypes	Mixed numerical + categorical	Hybrid (mean + mode)	Hybrid (Euclidean + matching)	Medium	Fast

`sklearn_extra.cluster.KMedoids`

```
class sklearn_extra.cluster.KMedoids(n_clusters=8, metric='euclidean', method='alternate',
init='heuristic', max_iter=300, random_state=None) [source]
```

k-medoids clustering.

Read more in the User Guide.

Parameters: `n_clusters` : int, optional, default: 8

The number of clusters to form as well as the number of medoids to generate.

`metric` : string, or callable, optional, default: 'euclidean'

What distance metric to use. See :func:metrics.pairwise_distances metric can be 'precomputed', the user must then feed the fit method with a precomputed kernel matrix and not the design matrix X.

`method` : {'alternate', 'pam'}, default: 'alternate'

Which algorithm to use. 'alternate' is faster while 'pam' is more accurate.

`init` : {'random', 'heuristic', 'k-medoids++', 'build'}, or array-like of shape

K prototypes

```
from kmodes.kprototypes import KPrototypes
import numpy as np

# Example data (replace with your actual data)
# Data should be a numpy array for input
data = np.array([
    [1, 'A', 2, 'X'],
    [1, 'B', 2, 'Y'],
    [2, 'B', 3, 'X'],
    [2, 'A', 3, 'Y'],
    [3, 'A', 4, 'X'],
    [3, 'B', 4, 'Y'],
])
# Specify the indices of the categorical columns
# In this example, columns at indices 1 and 3 are categorical
categorical_indices = [1, 3]

# Create a KPrototypes instance
# n_clusters is the number of clusters to form
kproto = KPrototypes(n_clusters=2, init='Cao', verbose=0)
```

k-prototypes minimizes a joint distortion function:

$$J = \sum_{i=1}^n \sum_{l=1}^k z_{il} \left[\underbrace{\|\mathbf{x}_i^{(num)} - \boldsymbol{\mu}_l\|^2}_{\text{numeric (k-means)}} + \gamma \underbrace{\sum_{j=1}^q \delta(x_{ij}^{(cat)}, m_{lj})}_{\text{categorical (k-modes)}} \right]$$

Where:

- $z_{il} \in \{0, 1\}$: assignment indicator
- $\boldsymbol{\mu}_l$: numeric centroid of cluster l
- m_{lj} : mode of categorical feature j in cluster l
- $\delta(a, b) = 0$ if $a = b$, else 1
- γ : trade-off parameter

K modes initialization

Huang Initialization → combinations created based on frequency

Cao → Density based → should be the default choice

Problem with K-means

- Convex shape
- Isotropic variance
- outliers

